

Paper

Simon van Hus

6147879

s.vanhus@students.uu.nl

Abstract

Abstract.

1 Background

1.1 Automatic Differentiation

Automatic Differentiation (AD), like the name suggests, involves programmatically finding the derivative of some programmed function [1]. The other main method for programmatically finding the derivative of a function is numerical differentiation, which uses the finite difference method. By adjusting the input(s) to the function by a very small number, we can see the effect on the output(s) of the function. Unfortunately, due to the way real numbers are represented using floating-point computation, this method is prone to round-off error (or truncation error). AD avoids this by actually performing the differentiation on a program, to produce the differentiated program. This is very similar to how a human would differentiate a mathematical function (sometimes called symbolic or manual differentiation), but performed on a computer program. AD makes very explicit use of the chain rule of partial derivatives of compound functions, which states that we can combine partial derivatives of parts of the function together into the complete derivative. For instance, when if some variable z depends on the variable y , which then depends on the variable x , we can obtain the derivative of x under z using the chain rule as in Equation 1.

$$\frac{dz}{dx} = \frac{\partial z}{\partial y} \cdot \frac{\partial y}{\partial x} \quad (1)$$

To actually implement automatic differentiation, we seek to break the target program down to its most basic mathematical operations. Then we can find the partial derivatives of these operations, and use the chain rule to combine them together into the derivative of the program. There are two main ways to actually resolves these derivatives: using either forward accumulation or backward/reverse accumulation. When applied in AD implementations these are commonly respectively referred to as forward-mode and reverse-mode. Both methods are described in the 1986 paper “The arithmetic of differentiation” by B. Rall [2].

In forward-mode AD we move through the program to differentiate in normal execution order. By knowing which input variable we wish to differentiate, we can compute every step of the derivative as our inputs are used by the program. Rall demonstrates this using a method known as dual-numbers, where each real number is represented by a pair of numbers, similar to complex numbers. In dual-numbers, the first number in the pair represents the primal part of the number, whereas the second number represents the derivative part (called the tangent in forward-mode). When we compute with these numbers through arithmetic operations, we can operate on the primal parts as normal, and use derivative rules to calculate the derivative of the result using the tangent parts. An example of this is given in Equation 2, where \dot{a} is the tangent part of some real number a .

$$(a, \dot{a}) \cdot (b, \dot{b}) = (a \cdot b, \dot{a} \cdot b + \dot{b} \cdot a) \quad (2)$$

Now we can find the derivative of some program with regards to the input x_i by setting x_i to 1, setting the tangents of all other inputs to 0, and just running through the program calculating tangents as we go. The tangent part of the output value(s) is also the calculated derivative of the whole program.

While forward-mode AD is fairly straightforward, it comes with some drawbacks. The main one being that for a function $f : \mathbb{R}^n \rightarrow \mathbb{R}^m$ with n inputs and m outputs, to get the effect of each input variable on each output variable, we would need to perform n passes over the function, one for each input variable (or we need to track n tangent parts for each step). This is can be cumbersome, especially if n is much larger than m . For those cases, we might be better off with reverse accumulation, or reverse-mode.

In reverse-mode, we peg the derivative part of one of our outputs with some seed (often 1), and set the derivative parts of the other outputs to 0. These derivative parts are generally referred to as adjoints instead of tangents in reverse-mode. When the outputs are set, we can work our way back through the function, calculating the derivative parts from the output to the input. Intuitively, this computes the change of the inputs for a small change in an output. Practically, the idea of working back through a program requires some way of knowing where the outputs came from (a sort of dependency structure). This then requires a forward pass, to find this structure, and often to setup any dual-numbers or other implementation details. And while reverse-mode is definitely harder to implement, it also provides us with a way to calculate the sensitivity of all inputs to an output, which is much more efficient for functions with many more inputs than outputs (which can be quite common in certain fields like neural networks).

In mathematical terms, calculating the partial derivative of one output with regards to one input, means calculating one cell in the Jacobian, the matrix of all partial derivatives. For a function $f : \mathbb{R}^n \rightarrow \mathbb{R}^m$ with n inputs and m outputs, the Jacobian would be a $n \times m$ matrix. Here a column i represents the partial derivatives $\frac{\partial \vec{f}}{\partial x_i}$, where \vec{f} are all outputs of f , and x_i represents a single input. A row j then represents the derivatives $\nabla f_j = \frac{\partial f_j}{\partial \vec{x}}$, where \vec{x} are all inputs of f , and ∇f_j is also known as the gradient of f_j . This is also shown in Equation 3, showing the Jacobian for some function f with n inputs (x_1, \dots, x_n) and m outputs (f_1, \dots, f_m) . An important take-away here is that forward-mode computes the derivatives of all outputs with regards to a single input, so a column in the Jacobian, and reverse-mode computes the derivatives of all inputs with regards to a single output, so a row in the Jacobian. Again, forward-mode is more efficient when we have more outputs than inputs or when the Jacobian has more columns than rows, and the reverse-mode is more efficient for functions with more inputs than outputs or for Jacobians with more rows than columns.

$$J_f = \left[\frac{\partial \vec{f}}{\partial x_1}, \dots, \frac{\partial \vec{f}}{\partial x_n} \right] = \begin{bmatrix} \nabla f_1 \\ \vdots \\ \nabla f_m \end{bmatrix} = \begin{bmatrix} \frac{\partial f_1}{\partial x_1} & \cdots & \frac{\partial f_1}{\partial x_n} \\ \vdots & \ddots & \vdots \\ \frac{\partial f_m}{\partial x_1} & \cdots & \frac{\partial f_m}{\partial x_n} \end{bmatrix} \quad (3)$$

While it was previously shown that reverse-mode automatic differentiation could be executed in time equal to some constant multiple of the execution time of the primal program [3], it seemed that a constant multiple of the execution memory was also needed, which could become very expensive for large programs. However, in 1992, Andreas Griewank showed that by using taping and checkpointing we could reduce the time and space complexity to a constant multiple of the log of the execution time [4]. In general the practice of taping refers to a form of tracing on the program we wish to differentiate, where we execute the program as normal and record all the steps and intermediate values in a first-in-last-out data structure referred to as a "tape" or Wengert list. In a second phase to the reverse-mode algorithm, the tape is then used to calculate the derivatives in question, which due to the first-in-last-out nature of the tape, is in the precise reverse of the execution order of the program. An important advantage of taping is that by giving each variable and intermediate calculation a unique id we can avoid redundant execution, because we can just refer to the intermediate value or tangent/adjoint stored in the tape. While taping is efficient time-wise, it clearly adds a memory overhead that can be quite sizable for large programs. Checkpointing aims to address this overhead by storing parts of the tape on disk, basically cutting the calculation of the derivatives in multiple steps. It should be noted that an input agnostic form of reverse-mode AD is also possible, where the primal program is transformed into a program that computes its derivative for any input, for which taping would only refer to storing the operations without any intermediate values. This is referred to as source transformation [5], and in its most basic form can be implemented as just interlacing the derivative calculations into the regular program. An example of this is provided in Listing 1, where we calculate the derivative of some variable y (as dy) with regards to the variable x_1 . For reverse-mode AD this kind of interlacing is not possible, as we need to reach the end of the program before we can start the reverse pass, which is exactly why we record our steps on the tape: so we can reverse over the tape and know how to produce our reverse AD program. An example of this is provided in Listing 2, where we again calculate the derivative of y (as dy), with regards to x_1 and x_2 .

```

x1 = 15
dx1 = 1
x2 = 7
dx2 = 0
r1 = x1 + x2
dr1 = dx1 + dx2
y = r1 × x2
dy = r1 × dx2 + dr1 × x2

```

Listing 1: An example of input-agnostic forward mode AD by source transformation, with the AD statements in red

For forward-mode AD, the evaluation of the derivative is done during execution. Like in 1996's FADBAD package, which provided both forward-mode and reverse-mode AD for C++ [6]. The reverse-mode uses the taping method described by Griewank, implemented through a method called operator overloading. In forward-mode, operator

```

x1 = 15
x2 = 7
r1 = x1 + x2
y  = r1 × x2

dy  = 1
dr1 = rdy × x2
dx2 = rdy × r1 + rd1 × 1
dx1 = rd1 × 1

```

Listing 2: An example of input-agnostic reverse mode AD by source transformation, with the AD statements in red

overloading refers to providing the basic mathematical operators with methods that work on the numbers represented by a pair (of a primal part and a tangent part); this is the dual-numbers approach we mentioned before. For reverse-mode, operator overloading is used to rewrite the basic mathematical operators, so they record their use and intermediate values to a single tape data structure. A similar implementation was also provided by Griewank et al. in the 1996 package ADOL-C [7], again in 2001 using more efficient expression templates by Aubert et al. [8], and later in 2014 by Robin Hogan [9].

Source-code transformation is eventually also implemented, in the Taped AD program [10]. Taped adds derivative calculations to the code, but also employs lazy/delayed evaluation in the forward pass. This allows Taped to do some activity analysis, which in turn allows it to combine or discard some partial derivatives to be more efficient. It also implements the previously discussed checkpointing, where part of the tape is stored to be restored and differentiated later. This, in theory, allows for differentiating programs of arbitrary size, because the differentiation process is not limited by the size of the working memory [11].

More recently implementations, like Fei Wang et al. 2019 paper, have shown how to simplify reverse automatic differentiation using continuation passing style and delimited continuation [12]. This method uses dual numbers and cleverly overloads operators so they call the forward pass as a continuation and then perform the backwards pass on the returned value. This can be efficiently implemented in functional languages, which is important because taping required either mutability or additional time complexity.

In 2022, Krawiec et al. show how reverse-mode AD can be extended efficiently to higher-order functional programs [13]. While the Wang paper also did this, Krawiec uses the functional nature to provide a correctness proof of the reverse-mode AD, something that had previously only been done on implementations that were either asymptotically inefficient or only worked on first-order languages. They do however need taping again to make it provable and efficient.

Vákár and Smeding provide a provably correct form of higher-order reverse AD without taping in their 2022 paper [14], based on earlier work by Elliott in 2018 [15].

And in 2022 as well, Schenck et al. show how to do both forward-mode and reverse-mode automatic differentiation on second-order array language with nested data parallelism [16]. They do this by eliminating taping again, which forces sequential execution, by allowing potential redundant execution. But by limiting their AD implementation to second-order functional languages, they can largely avoid this redundancy with efficient

program transformations on parallel operators.

Finally, in 2023, Smeding and Vákár bring back explicit dual-numbers to reverse AD [17]. However, instead of pairing each number with its computed adjoint, they instead pair it with a linear backpropagator function, which they can then later chain to get the full derivative. While this initially seems to eliminate the need for taping, they find that through optimizations they return to a concept that is very close to taping and show that it is in fact equivalent.

1.2 Tracing

Tracing is a concept in computer science that is often left without proper definition. While the main ideas behind tracing are well known, they are generally assumed known by the reader and therefore left without explanation. This is also in part because, in software engineering, the term tracing also refers to finding the origin of some call (“tracing” the call stack), which is tangentially related to the tracing we are interested in, but can leave definitions of tracing a bit muddled. This is why, in Section 2, we will discuss more about that proper definition. For now, it is important to know that, when we refer to tracing in this paper, we speak about tracing the path of computation through a program, given some (valid) input to said program. In other terms, given a program and an input, we walk through the program and record each computational step for some later purpose. This recording can happen with some domain-specific pseudo-language, or in full fledged code if we wish to reevaluate the trace later on (or a combination of the two).

Doing tracing gives us some interesting insights into a program we trace. First, it effectively ignores control flow. This is fairly intuitive, when given a set of inputs to a program, the control flow will control what path the program uses, and since we only record computations we find what is often dubbed a “straight-line program” for some inputs. This can be useful for instance in Automatic Differentiation, where we often only want to differentiate a computation, not the entire program including unused branches. This will also be the use of tracing in this paper, as laid out in Section 3.

Tracing also often allows us to simplify the program in more ways than removing control flow. This is through type selection, where we only select a subset of the original types present in the original program to be used in the trace. For composite types, types constructed of other types, and collections, this is somewhat obvious: we can rewrite a pair of two values, into two different variables, which removes the pair from our type system in the trace. The same can also be done for arrays, something that will prove useful for the tracing and differentiation of functional parallel array programming, described in Sections 1.3 and 2. For some instances it also allows us to remove non-composite types from the type system in the trace. For instance, as we remove control flow, we might want to ignore Boolean values altogether, especially if our input and output do not include any Boolean values. This is highly dependent on which information we wish to see in our trace exactly, but if Boolean values are never actually used, they may as well be ignored. In this case, we can “trace away” the types we wish to ignore by omitting any values of this type or any computational steps that produce this type. However, it should be noted that this is not always possible, for instance if a type τ that we wish to trace away is used to construct a type σ we wish to keep; tracing away the type τ would also mean tracing away the construction of the type σ leaving us with an incomplete trace. Of course, perhaps we can rewrite the production of the type σ from τ in our

trace, so we do not need τ , however this would need to be done on a case-to-case basis for different methods $\tau \rightarrow \sigma$. Regardless, tracing away types or keeping them, allows us control over what information we keep in the trace, and what information we choose to ignore.

As mentioned, in literature we see this type of tracing used for automatic differentiation. One of these uses was by Bischof in 1991 [18]. In his paper, Bischof discusses the use of the computational graph of a program in automatic differentiation (using ADOL-C [7]). The computational graph of a program is an directed acyclic graph, where each node contains a computational step in the program, and edges connect these steps in the execution order of the program. Bischof creates this graph from the tape produced by ADOL-C, which makes sense: for automatic differentiation as discussed, the tape acts as a sort of trace, recording the steps that are important in the automatic differentiation. Bischof then uses a graph colouring algorithm on the computational graph to highlight “component functions” that may be differentiated concurrently, as to improve the running time of the algorithm. In 2008, Bischof et al. expand on this by extending the tracing automatic differentiation to loops [19]. They do this by extending ADOL-C, paying specific attention the parallelization opportunities present in automatic differentiation.

In a similar vein, Dougal Maclaurin presented in his PhD thesis in 2016 [20] a paper introducing Autograd. A software package to automatically differentiate Python code (including AD for the vector library Numpy). As Python is an expressive JIT-compiled (Just In Time) dynamic language, they opt for tracing to construct the computing graph on the fly when a function is called, and like Bischof’s work this allows them to do the backwards pass off reverse-mode AD on the computational graph. They do this by wrapping their variables as “nodes” in the computational graph. When a variable is used, it is first unwrapped for use, and then the result of whatever operation used the variable, is stored as a new variable and wrapped as well. The original variable and the produced variable are then linked such that the produced variable stores a reference to the original variable. This creates the reverse computational graph, which is exactly what is needed for the reverse AD pass.

TensorFlow, a machine learning library, also uses tracing to create computational graphs [21]. This kind of tracing is not as low to the ground as actually following individual computations. Since TensorFlow mainly focusses on building artificial neural networks, the computational graph is made explicit by the programmer. While there are some nuance differences between a computational graph of a neural network and the neural network itself, these differences are somewhat unimportant. More interestingly, TensorFlow allows for the partial execution of the computational graph. While Bischof’s use of a graph colouring algorithm already suggested this, TensorFlow actively uses this technique to re-run partial computational graphs, which works well for the explicit nature of neural networks, as the computational graph stays unchanged even if the inputs change (as neural networks do not have internal control flow).

Finally, 2018’s JAX uses tracing to enhance performance of general machine learning code [22]. The programmer annotates functions to be analysed by JAX, which then traces as optimizes them. Rather than finding the computational graph (or predefining it), JAX waits for Python to execute the function and actually traces it. Then, JAX optimizes it, mainly through a process called fusion, which is discussed in Section 1.3. This is also where JAX gets its name: Just After eXecution, as it waits for Python to execute the function first. It should be noted that JAX can only do this for functions which are pure-and-statically-composed (PSC), meaning functions that have no side-

effects and that do not change with different inputs. Again, machine learning code is especially suited for this, as it often already satisfies this PSC assumption.

1.3 Functional Parallel Array Programming

Array programming languages are programming languages that treat the array as a central data structure. This generally includes that functions, both user-defined and built-in, could be applied to arrays through vectorization. Vectorization involves applying a function to every element of an array at the same time. For instance, vectorization of addition would add two arrays together element-wise. This is shown in Equation 4, where \vec{a} and \vec{b} both are arrays of the same size.

$$\vec{a} + \vec{b} = [a_1 + b_1, \dots, a_n + b_n] \text{ iff } |\vec{a}| = |\vec{b}| \quad (4)$$

In general vectorization would only work for arrays of the same size were it not for another central concept: broadcasting. Broadcasting involves the resizing of arguments to functions so they can be used. A very clear example would be if we wished to add a scalar value to each element in an array of scalars. To do this with vectorization alone would mean we'd need another array which replicates the scalar we wish to add for each element in the array we wish to add it to. Broadcasting basically does this for us, as exemplified in Equation 5.

$$\vec{a} + 2 = [a_1 + 2, \dots, a_n + 2] \quad (5)$$

Array programming languages also often support higher-order operators for use on arrays. An important operator for arrays is fold (or reduce), which applies a binary function to elements in an array, where one argument accumulates the previous results. It is easy to imagine how such an operator could be used to, for instance, sum all the items in a 1-dimensional array. An important realization is that, since fold only returns the final result, fold can reduce the dimensions of an array by one. In our summation example, we fold a one-dimensional array into a zero-dimensional array, namely a scalar value. Similar to fold is scan, which like fold applies a cumulative binary function to each element in the array, but rather than returning only the result, it returns all intermediate results in an array (with the last element being the final result).

Other important array functions include map, which applies a function to each element in an array. Then, forward permutation (scatter) and backwards permutation (gather), which permute one array into a new one by respectively mapping the indices of the source array to those of the new array or the indices of the new array to those of the source array. Generate, which generates a new array as well, but does by taking the dimensions of the desired array and a function that takes in an index and outputs a value.

We should also not gloss over the actual implementation of these arrays, especially in functional languages where there exists two major ways of constructing arrays [23]. Pull-arrays are the more used of the two, here arrays are represented with a function from an index to a value. In push-arrays, consumers are provided with a method to write into memory. This means that the way that the efficiency of array operations can change based on the array representation. For instance, indexing is faster on pull-arrays while push-arrays are quicker to concatenate. This basically divides the array operations in two camps: push-operations and pull-operations.

These two camps play an important role in a concept of fusion. When we have multiple back-to-back parallel array operations, executing them naively introduces a lot of overhead for reading and writing intermediate values to memory. Instead fusion allows us to combine these operations together, so we can compute them in one go without the overhead of storing intermediates. However, we cannot just go chaining parallel operations, not all parallel operations fuse together nicely. In fact, pull-array operations only fuse with other pull-array operations, and the same goes for push-array operations. This means for instance that we can fuse multiple scatter operations, but not a scatter and gather operation.

Now the reason for choosing a (functional) array language over a general language is often because we need to process large amounts of numerical data, and arrays are well-suited for parallelism. To be precise, we are talking about data parallelism here. Task parallelism is when two or more computer processes run simultaneously on different processor cores. Data parallelism is when an operation (or a string of operations) is done element-wise on data structure like an array. The parallelism of data parallel processing of these operations on each element, rather than the parallelism of different processes. An important distinction between task and data parallelism, is that while parallel threads in task parallelism can generally start, run, and end independently of each other, data parallelism threads move in lockstep with each other. This is lockstep or synchronous execution means that the execution does not continue until the current operation has been applied to all elements in the array, which may be important if we want to do multiple parallel operations back-to-back. Furthermore, modern GPU architectures are especially well-suited for this type of synchronous parallelism, as graphics processing overlaps in large part with parallel array processing.

A good starting point for the history of functional parallel array programming was in 1992, with G. Belloch's paper on the parallel array programming language NESL [24]. The language was strongly-typed and had no support for side-effects, making it a functional language. The main way to add parallelism was through the inherently data-parallel "vectors" the language introduces in lieu of lists. These vectors could also be nested, and functions could run in nested parallel on these vectors. Another major inclusion was to allow user-defined functions to be run (in parallel) on these vectors, making it possible to write more complex nested data-parallel algorithms than before.

The functional language Haskell, saw the introduction of task-parallelism well before its first official release, through libraries like pH [25]. Some data-parallelism followed [26, 27, 28], but this was limited to applying a function over a flat array. However, in 2001 nested data-parallelism was introduced to Haskell by the NEPAL project by Chakravarty et al. [29]. The paper largely focusses on reimplementing NESL as a Haskell library, but creates a much more expressive data-parallel language doing so. This is because NESL was rather limited in scope, whereas Haskell was already a fully-fledged functional programming language. Two important concepts come to the forefront in the NEPAL paper, namely flattening and fusion. Both in NESL and in NEPAL, higher-dimensional nested parallelism is "flattened" to a single distributed parallel operation. In NESL, this meant that data-types had to be limited to tuples and the vectors it introduced, to make sure this flattening operation worked correctly. Since then however, Keller and Chakravarty had shown this flattening transformation could also be applied more generally to cover the full range of types of general programming languages [30, 31, 32]. This allowed them to apply the nested data parallelism of NESL to a more expressive language Haskell with NEPAL. Furthermore, they also showed that in combination with fusion it could produce efficient code for distributed machines [33]. Fusion is where multiple separate parallel operations are combined into a single parallel operation,

which greatly improves performance of complicated parallel programs. This is important because many operations on arrays introduce the need for intermediate arrays to be computed. Doing this in parallel leads to more problems, as these implementations rely on gang parallelism, where the parallel threads remain in lockstep with each other [34]. Fusion helps us here, as we can reduce the number of intermediate arrays to be generated, as we can calculate the results of multiple operations at once [35, 36].

All this work culminated in 2007’s Data Parallel Haskell (DPH) [37], by Peyton-Jones et al. Its main feature was the parallel array, that like NESL’s vectors, was the main way of adding parallelism to a program. However, these parallel arrays could now hold any type, such as other arrays or functions, like Haskell’s native (non-parallel) lists. Furthermore, DPH provides parallel variants of Haskell’s native list functions, and a parallel alternative to Haskell’s list comprehensions. The main difference between Haskell’s native lists and DPH’s parallel arrays (besides the parallelism) was that evaluating any value in a parallel array would require evaluation on all the array’s elements, whereas Haskell as a lazy language would not normally do that. This is to be expected, as parallelism becomes meaningless if it is only applied to a single entry of an array.

Outside of Haskell, a functional array-programming dialect of C was developed: Single Assignment C (SAC) [38, 39, 40]. It would go on to distinguish itself as a functional array programming language in a style more familiar to programmers of imperative languages (like C). The main mechanic in SAC is the with-loop, which takes a generator that dictates a looping mechanism and an operation that dictates the return value. These operations can be functions like “fold” to reduce the rank of an array, or “genarray” to generate new (multidimensional) arrays. Besides the imperative style, the main draw of SAC is that its performance is comparable to Fortran and C, while its programs are generally more concise (for intensive numerical applications.)

In 2010, Keller, Chakravarty, et al. presented a new data-parallelism approach for Haskell in “Regular, Shape-polymorphic, Parallel Arrays in Haskell” [41]. Previous approaches had focussed on irregular arrays, where an array could contain arrays of different lengths. The library Repa, introduced in this paper, was made for regular arrays where arrays of each nested rank are the same size. However, this allows the library to be purely functional and support shape polymorphism. While DPH was purely functional as well, it wasn’t especially performant on regular arrays and it also did not support shape polymorphism. In shape polymorphism, the type of a collection is fixed (unlike in type polymorphism), but the shape of the collection is not [42]. For instance, under shape polymorphism a function may be applied to either a flat array, or a 10-dimensional one. While shape polymorphism for functional arrays had been implemented before in SAC, Repa implemented it by embedding it into Haskell’s type system, whereas the SAC implementation had required a purpose-built compiler. This also allowed programmers to more easily see and control the shapes of their multidimensional parallel arrays, and build their own shape polymorphic parallel functions.

In 2011, Repa was succeeded by the Accelerate project [43]. Accelerate is a library for Haskell, aimed specifically at bringing parallel array programming to modern GPUs. It mimicked many of Haskell’s native list functions with parallel alternatives (that run on the GPU), and used the typed shaped polymorphism from Repa. It also separated “collective” (array) computations and scalar computation by wrapping these in Haskell monads. Here, collective computations could include scalar computations, but not the other way around. This meant excluding nested and irregular data parallelism, which in turn allows Accelerate to efficiently run on GPUs (which are much more constrained than CPUs). It also meant that these arrays could only contain scalars, no functions or

other types.

Another interesting example of a parallel array programming language is Remora by Slepak et al. [44]. The language, inspired by earlier array programming languages APL [45] and J implements rank-polymorphism. Rank polymorphism is similar to shape polymorphism, but it annotates functions and operators with an array rank they can operate on, and was also present in Repa and Accelerate. Remember that scalars are considered rank 0 arrays, a flat array is rank 1, a matrix is rank 2, et cetera. In rank polymorphism, arguments are transformed (re-ranked) such that they are the rank required for a specific function or operator. Specifically, an operator defined for a certain rank, is automatically defined for any higher rank, because it can be mapped over these higher dimensions. This is subtly different from the more general shape-polymorphism, as rank only refers to the number of dimensions, while shape also contains information on the size of these dimensions. With Remora, Slepak et al. tried to shed some light on the more “murkier corners” of the array-computational model. They do this by generalizing the array-computational model, which then allows them to both address some of the shortcomings of APL, but also allows them to extend the model to allow arrays of functions and arrays of arguments, which in turn allows for the parallel MIMD (multiple instruction, multiple data) architecture, rather than only SIMD (single instruction, multiple data) parallelism.

In 2017, we got one of the major current functional data-parallel array languages in Futhark [46]. Futhark’s design focusses on efficient nested data-parallelism. They do this by using both “aggressive” fusion (fusing as much as possible), followed by flattening (like we saw in NESL). Finally through some more optimizations, Futhark produces very performant programs. To facilitate this performance however, they do not support higher-order programming, as Futhark only supports up to second-order.

Finally, a more recent parallel array programming language is Dex [47, 48]. Rather than avoiding loops and explicit indexing, like NESL, NEPAL, DPH, and Repa had all done, Dex suggests that these features might introduce more clarity, if only they were implemented correctly. The main idea is to treat index sets as types and arrays as functions. In reality this “index comprehension” can also be seen as functions that return arrays, and allow declaring iteration over multiple dimensions in a single line. Of course, this is the same idea as pull-arrays, a representation also used by Accelerate under the hood. However, the main novelty of Dex is that they use this to make explicit loops, which in turn makes some parallelism opportunities also explicit. Also when these index comprehensions are presented back-to-back, opportunities for fusion become fairly clear as well. In their paper, they also show that on some benchmark problems, Dex performs similarly to Futhark, as a functional array programming language that was specifically designed to write performant parallel GPU code.

2 Tracing

In the broadest terms, when we trace a program, we track the most basic steps the program takes provided some input. This is relevant for many applications in Computer Science. For example, certain automatic differentiation (AD) effectively implement the forward-pass as tracing, and then perform the reverse pass on the trace^[cite me]. Tracing is also used in Artificial Intelligence, where tracing applications can help determine how much memory needs to be allocated, which can speed up training if the model is run multiple times^[cite me].

However, despite its ambivalence, tracing is rarely properly defined, or defined only for a specific use case. So, in this section we set out to create a more general definition of tracing.

To start, it will help us along to set clear expectations for what we expect a tracing function to do. In the simplest terms, we expect a tracing program to take an input program with a set of inputs, and output a “trace”. This output trace is defined as a sequence of operations the input program performed on the inputs to get the expected output. A term often used for an output trace is a “single-line program”^[cite me]: a program without control flow. Clearing control flow like if-then-else statements is only natural: after all, provided some input the program will only walk down one variation of this branching path.

Furthermore, it is also generally accepted that the trace consists of a subset of the syntax of the input program. Because we are generally more interested in what happens to the data in our program, we can “trace away” functions and data structures. More precisely, say our input program has the types as defined in Equation 6, where we have sum-types as $\tau + \sigma$, product types as $\tau \times \sigma$, functions as $\tau \rightarrow \sigma$, literal real numbers, and literal Booleans.

$$\tau, \sigma := \tau + \sigma \mid \tau \times \sigma \mid \tau \rightarrow \sigma \mid \mathbb{R} \mid \mathbb{B} \quad (6)$$

We can imagine our simplified language, in which we will express our trace – as a language with fewer type formers. By choosing a subset of the type formers in our program, we can indicate which data structures should be traced away. A common option is to keep only “ground types”, where we defined a ground type as a type that is not constructed of other types. Looking at our example in Equation 6, a trace keeping only these ground types would keep only the real numbers and the Booleans as they are not built of other types. Another common option is to keep only continuous types, tracing away all unground and discrete types. Doing that on our type set in Equation 6 would leave us with only the real numbers. This is under the assumption that the discrete types are not actually used as data we are interested in tracing of course, but since tracing will remove all control flow from the program, keeping Booleans and operations on Booleans intact may be meaningless.

The main take-away here is that there is some freedom of choice in what to trace away. What parts we keep and what parts we trace away is very dependent on what information we want to keep in our trace, which in turn is dependent on what our exact goal is for the tracing in the first place.

We can also choose to keep some of our unground types, but then we run into a problem. Say we keep only functions ($\tau \rightarrow \sigma$) and real numbers, but our input program contains a function with type $\tau \rightarrow (\sigma_1 + \sigma_2)$. This typing is valid in our input program, but no longer valid in our trace, so we find ourselves in a bind. It will be impossible to trace away the sum-type in the output of the function without tracing away the function itself. This is because tracing something away basically means either deconstructing or ignoring it in the trace. For instance, tracing away a tuple, would mean tracing the individual components of that tuple to trace it away. Whereas keeping things in the trace means just keeping them untouched. Therefor, we cannot keep a type like a function $\tau \rightarrow (\sigma_1 + \sigma_2)$ in our trace, because we cannot access the sum type without tracing away the function. Of course we could define a subset $\tau', \sigma' := \mathbb{R}$ and then redefine (or add a definition for) our function so that it becomes $\tau' \rightarrow \sigma'$ making it safe

to trace. This then underlines the rule at work here: we can only keep types that do not be constructed of types that are traced away. This is why the ground types are a natural set of types to keep, as they are never constructed from other types.

In a similar vein, we may also encounter operators in our trace that take in or produce types that are not allowed in our trace. For operators that produce a type that is not in our trace, tracing them away is no problem. Since we know we will not be interested in whatever output they produce for our trace, we can simply omit them from the trace altogether. For instance, if we keep only real numbers in our trace like before, an operator returning a Boolean value is of no interest for the trace. However, this is not a simple for operations that take in a type we wish to trace away, yet produce a type we wish to keep in our trace. A simple example of this is the “switch” operator, which takes in a Boolean value and two values of another type, of which it returns one depending on the Boolean value (see Equation 7).

$$\begin{aligned}\text{switch}(\top, a, b) &= a \\ \text{switch}(\perp, a, b) &= b\end{aligned}\tag{7}$$

While the switch operator looks like it mimics if-then-else statements, it is generally accepted that it does so in a non-lazy way^[cite me], where both a and b are evaluated before returning either. The main problem here is that we wish to keep operators that produce types we keep in our trace, yet we do not wish (or are not even able to) express the Boolean value in our trace. Now, due to switch statement’s likeness to if-then-else, the solution here is pretty clear: only trace the value that gets returned. However, it is not always that easy: as we introduce arrays and array operations in Section 2.4, we will see how operations like mapping on an array need a special solution.

This all is to say that the while we can either ignore or homomorphically copy basic operations for our trace, sometimes we need a special solution. This is mainly because we do not want to lose the information that is needed to execute the trace as a single-line program, even if that means fudging our operations a little. This also means that, while the operations in our trace language might be a subset of the operations in the original expression language, they might contain modified operations

It seems that our tracing definition comes down to a function that takes in a program and an input to that program, and outputs the steps taken by the program run on the input. Where the input program takes uses some set of types, of which only a subset is kept in the trace, where the types in this subset may not be constructed using types from outside of the subset. What now remains is a concrete definition of the output of the tracing program. We have already stated that it should somehow contain the steps done in by the input program. The steps we wish to record are generally basic operations like arithmetic operations. But other operations, such as operations on arrays, can also be added depending on the ultimate goal of the tracing. More importantly, as we expect our trace to be akin to a single-line program, we may consider our trace as a series of let-bindings, akin to A-normal form^[cite me]. This means storing each operation as a pair of a unique name or id and the operation performed (like the name and value of the declarations in a let-binding).

2.1 Tracing Correctness

Before going into specifics on how to implement tracing, it would also be a good idea to formalize when a trace is actually correct. Like we posed before, we start with some

program formed from some expression language S , and some input I that is valid for that program. If we would wish to resolve a program S on input I , then we would need some evaluation function that produces the expected output O . Now, given some trace language T we can write a tracing function that gives us the trace and output of a specific program and input combination. We can write this out as the two functions `eval` and `trace` in Equation 8.

$$\begin{aligned} \text{eval} &: S \times I \rightarrow O \\ \text{trace} &: S \times I \rightarrow T \times O \end{aligned} \tag{8}$$

With this we can formalize two criteria for our trace. First, the trace, as a single line program $t \in T$ produced by the trace function needs to produce the correct output. Now, as mentioned before, t might contain transformed operations, that are not present in S . Therefor we either need to look at traces $t \in S \cap T$, or use a different evaluation function. For now we will use the former, to assert the output criterium in Equation 9. Here we state that for any program s with any input i : if the trace t is also a valid program in S , that the evaluation of t on i should be the same as the evaluation of s on i or the output o we got out of the tracing function.

$$\begin{aligned} \forall s \in S \\ \forall i \in I \\ \text{trace}(s, i) &= (t \in T, o \in O) \\ (t \in S \cap T) &\rightarrow (\text{eval}(s, i) = \text{eval}(t, i) = o) \end{aligned} \tag{9}$$

Furthermore, tracing a trace t should also return that trace t . This is because we want to find the minimal straight-line program using tracing, and if tracing the trace we found reduces it somehow to a more minimal program, we know that the original trace was incomplete. This is expressed in Equation 10, where we assert that for some program $s \in S$ and some input $i \in I$, the trace t (produced by tracing s on i), is the same as the trace obtained from tracing t itself.

$$\begin{aligned} \forall s \in S \\ \forall i \in I \\ \text{trace}(s, i) &= (t \in T, o \in O) \\ \text{trace}(t, i) &= (t, o) \end{aligned} \tag{10}$$

The above statements, assert that a trace should produce the correct output value as expected from the input program, and that a trace should be its own trace. While these assertions do not say a lot about the nature of the actual trace, they do set some baseline requirements for the trace, and proving the correctness of a trace. This vagueness on the contents of the trace is partly because we cannot really say anything about a trace without dissecting the source program as well, which would bring us to a point very close to actual tracing itself. In another part however, this is because we do not want to make any assumptions what can or cannot be in our trace. While it is likely that some there is significant overlap between S and T , as mentioned, we might need some additions to T to actually be able to trace everything in S correctly. Also, whilst in practice it might be meaningless, a trace where $T = \emptyset$ is in itself not incorrect: any trace would simply be empty. In a similar vein a trace where $S \subseteq T$ would also be meaningless in practice, it is also not wrong: any trace would simply be the same as the source program.

As an additional note, Equation 9 also implies something interesting. If we want our trace to output the same value as the original program, we cannot trace away the type of the original programs output. Say we trace away Boolean values when we are tracing a program that returns a Boolean value, then we find ourselves stuck, because we trace away all operations that produce Boolean values. And of course, if our trace is not allowed to produce any Boolean values, we cannot produce the required output either. Therefor we must assure that the type of the output is valid in our trace as well.

2.2 Basic Tracing

We now define some basic tracing steps for some arbitrary language. For clarity's sake, we will do this is with Haskell^[cite me]code. To do this we first define a language and values on which we will operate. We do this in Listing 3, where we define a basic lambda calculus. Here the value types are represented as the algebraic data type (ADT) `Value`, where we find constructors for Booleans (`VBool`), real numbers (`VReal`), and functions (`VFunc`). Then we define the four terms of a basic lambda calculus in the `Expression` ADT: application (`EApply`), abstraction (`ELambda`), loose values (`ELift`), and variable reference (`ERef`). To make tracing a little more interesting we also add in if-then-else statements (`EIf`) and binary operators (`EOp2`). For those binary operators, we define four operations in the separate `Op2` ADT: addition (`Add`), equality (`Equ`), multiplication (`Mul`), and inequality (`Neq`). Finally, to make use of variable references, we define an environment as a mapping of strings to values. We interact with this environment in two ways: by inserting values into them, and getting values from them (indexing). The function signatures for these interactions, respectively `insert` and `(!)`, have been included in Listing 3 as well. We can use this language and evaluate it, an example of this has been provided in Appendix A.

```

1 data Value = VBool Bool | VReal Float | VFunc (Value -> Value)
2
3 data Expression
4   = EApply Expression Expression
5   | EIf Expression Expression Expression
6   | ELambda String Expression
7   | ELift Value
8   | EOp2 Op2 Expression Expression
9   | ERef String
10
11 data Op2 = Add | Equ | Mul | Neq
12
13 type Environment = Map String Value
14
15 -- Operations on maps:
16 -- (where Map a b is a mapping from keys of type a to values of type b)
17 insert :: a -> b -> Map a b -> Map a b
18 (!) :: Map a b -> a -> b

```

Listing 3: Minimal lambda calculus with added if-then-else and binary operators

With our language in Listing 3, we can almost start tracing. However, we must first decide which parts of the language we keep, and which parts we wish to trace away. In the previous section, we talked about how we can do this by selecting which type formers we wish to keep. In Listing 3, we have practically defined the types of our values by the

data constructors present in the `Value` ADT as Booleans, real numbers, and functions. Let us now choose to keep only real numbers in the trace.

We now define a new ADT for traced values in Listing 4. This is only so we can incorporate a name into the values we wish to keep in our trace. These names will help us read the trace, and can be incrementing numbers or something entirely random, as long as they are unique. The basic idea is here to feed the `trace` function a number with which to generate the steps' names from, and increment the number every time we do. However, since this clutters the code while not being very interesting, we will assume we have some function `getName` that provides us with a unique name. Furthermore, it is important to see that we still have Boolean values and functions in our `TValue` ADT, even though we only wish to keep real numbers in our trace. This is because we might still need these values to resolve expressions, even if they never end up in the trace. We might also achieve this by extending our original `Value` ADT (from Listing 3) with traced variants of values, but this is merely a point of preference. Finally, we have also changed the signature of the function value to return a trace as well, as we move on to functions we will see how this works.

```

1 data TValue = TBool Bool
2             | TReal String Float
3             | TFunc (TValue -> (TValue, Trace))
4
5 data Traced = TLift TValue | TOp2 Op2 String String
6
7 type TEnvironment = Map String TValue
8
9 type Trace = [(String, Traced)]
10
11 getName :: String

```

Listing 4: Basic trace building blocks

First however, with our basic building blocks for tracing set up, let's trace away these boolean values. We do this with the trace function in Listing 5. For now, we will leave out abstraction and application, as it might be easier to talk about tracing away Booleans first.

When we trace away Booleans, like in Listing 5, it is useful to think about where these Boolean values actually come up. In our minimal language from Listing 3, there are only three points: when they are included as literal values, as the input or output to basic operations, or as the conditional in if-then-else statements.

Let us start with the easiest first: literal Boolean values. When we encounter literal values during tracing, and they are of a type we wish to keep for our trace, we simply add their instantiation to the trace (as `TLift` in Listings 4 and 5). This is extremely straightforward: those values might be used by the operations we wish to trace, so they should be included in the trace themselves as well. For values of types we wish to trace away, we simply do not include them in the trace. After all, our trace should be fine without them, as we do not include any operations that require them in our trace, right? For now this seems obvious: if we look at the language in Listing 3, we see that there are no other uses for Boolean values than the use in the equality and inequality operators, and as the conditional in the if-then-else statement. Since we plan to trace these away, we do not appear to need these value instantiations in our trace either. However, at


```

1 trace :: TEnvironment -> Expression -> (TValue, Trace)
2 trace n (EIf e1 e2 e3) =
3   -- Since we e1 should resolve in a Boolean value, we do not need to trace it.
4   let v1 = eval n e1
5   in case v1 of
6     -- We can check for the type of v1 and its value in one go
7     -- We trace only the relevant branch
8     (VBool True)  -> trace n e2
9     (VBool False) -> trace n e3
10    -              -> error "Type mismatch in trace/EIf"
11
12 trace n (ELift v) =
13   -- Check if v is a value we would like to trace
14   case v of
15     -- If yes return the transformed value with its simple trace
16     (VReal v) ->
17       -- Generate a name for this step and make the TValue
18       let s = getName
19           v' = TReal s v
20       -- Combine the TValue with a trace of its instantiation
21       in (v', [(s, v')])
22     -- If we do not wish to trace something, we can just return the value
23     -- with an empty trace.
24     (VBool v) -> (TBool v, [])
25     -- Instantiation is not allowed for functions, they need to be
26     -- abstracted using ELambda
27     -              -> error "Type mismatch in trace/ELift"
28
29 trace n (EOp2 op e1 e2) =
30   -- We again first trace e1 and e2
31   let (v1, t1) = trace n e1
32       (v2, t2) = trace n e2
33   -- We get a ready name in case we need it
34   s = getName
35   -- This case syntax allows us to select for the right operator with the
36   -- right value types at the same time.
37   in case (op, v1, v2) of
38     -- Since add and mul take in reals and produce one too, we trace both
39     -- the operation and the origins of v1 and v2
40     (Add, TReal s1 a, TReal s2 b) -> (TFloat s (a + b),
41                                       (TOp2 op s1 s2) : t1 ++ t2)
42     (Mul, TReal s1 a, TReal s2 b) -> (TFloat s (a * b),
43                                       (TOp2 op s1 s2) : t1 ++ t2)
44     -- For operations producing Booleans we only return the result, but they
45     -- are not traced, and therefore return an empty trace
46     (Equ, TBool _ a, TBool _ b) -> (TBool (a == b), [])
47     (Equ, TReal _ a, TReal _ b) -> (TBool (a == b), [])
48     (Neq, TBool _ a, TBool _ b) -> (TBool (a /= b), [])
49     (Neq, TReal _ a, TReal _ b) -> (TBool (a /= b), [])
50     -              -> error "Type mismatch in trace/EOp2"
51
52   -- There is nothing to trace when fetching a variable, but we still need to
53   -- actually get the value
54   trace n (ERef s1) = (n ! s1, [])

```

Listing 5: Tracing away Boolean values

the end of Section 2.1 we already discussed what would happen if our program were to return a Boolean value. And in Section 2.4, we will see how this might not be entirely true when we talk about arrays and operations on arrays like mapping a function.

Tracing (away) simple operations like addition and equality (**EOp2**) are done in a similar vein. If the operation returns a value of a type we wish to keep in our trace, we include the operation in our trace as well. Similarly, if the operation returns a value that we do not wish to keep, we simply do not trace it. Again, if there was an operation that took in a value of a type we do not wish to trace, and returned one that we do wish to trace we run into a problem. Luckily, these operations are not included in our current example.

When we trace an if-then-else statement, we know we have to deal with a Boolean regardless. Luckily for us, we know we only need to trace one of the branches. This means quite simply, that we can ignore the if-then-else statement, and act like the program continued at the branch that is chosen. Since the input is provided, we can resolve the conditional immediately, and then just trace the appropriate branch.

Finally, tracing variable references are simple as well. Currently the only named variables that occur are those created in lambda abstractions or those that are provided as inputs. But no matter how they are created, variable reference does not require tracing. This is because the trace will reference the values regardless of whether they are instantiated on the spot or somewhere previously. And if they were defined previously, that definition is already in the trace somewhere.

2.3 Function Tracing

With our basic tracing established, we can now talk about tracing functions, which are more complicated. It is the tracing of abstracted functions that is the first issue here. The issue is that when we perform an abstraction (as with **ELambda**), there is nothing to trace. In fact, we can see this as an instantiation of a function literal, and when functions are not in our set of types to keep in the trace, this abstraction creates an empty trace. However, leaving it at that would mean we never actually trace the body of the function. Yet, at the time of the abstraction, we also do not yet know the input to the function either, meaning we cannot trace the body at that time. We must instead consider how we delay tracing until the function is actually applied. This is where our notation for **TFunc** (as in Listing 4) comes up. We wish that functions while tracing perform tracing themselves, thus return a **Trace** together with the return value. This is then what we do in the abstraction step: we set the trace on the body of the function as the body of the function we return. Similarly, we also give this tracing function call the environment at the time of abstraction, allowing the function body to access any free variables that were defined at that time. This makes application also very simple: we apply the function, and then just combine the trace of the functions instantiation, with that of the argument, and that of the functions execution. We also trace the functions instantiation, since at the time of application we do not now if the expression that leads to the function does anything else that we might need to trace as well. Finally, this is results in what we see in Listing 6, where we left out any patterns of trace that were already present in Listing 5.

```

1 trace :: TEnvironment -> Expression -> (Value, Trace)
2 trace n (EApply e1 e2) =
3   -- First trace e1 and e2
4   let (v1, t1) = trace n e1
5       (v2, t2) = trace n e2
6   -- Check if v1 actually returns a function
7   in case v1 of
8     -- Do the application, return the result and the combined trace
9     TFunc f -> let (vf, tf) = f v2
10                in (vf, tf ++ t2 ++ t1)
11    _        -> error "Type mismatch in trace/EApply"
12
13 trace n (ELambda s e1) =
14   -- Define the function, insert value x as variable s into the environment that is currently
15   -- present, and trace the body
16   let f = TFunc (\x -> trace (insert s x) e1)
17   -- Return the function as abstracted function as a value, and no trace
18   in (f, [])

```

Listing 6: Tracing away functions

2.3.1 Tracing let bindings

As an additional structure present in functional languages that we might wish to trace, there are let-bindings. Recall that let-bindings are effectively the same as lambda abstractions that are resolved immediately. This makes it extremely easy to resolve them, because we can just trace the let-side of the binding and add it to the environment for the tracing of the righthand-side.

Adding let-bindings and tracing them is done in Listing 7.

```

1 data Expression
2   = ...
3   -- The string here is the name of the bound variable
4   | ELet String Expression Expression
5
6 trace :: Environment -> Expression -> (TValue, Trace)
7 trace n (ELet s1 e1 e2) =
8   -- Evaluate e1 first, then e2 with e1 in its environment
9   let (v1, t1) = trace n e1
10      (v2, t2) = trace (insert s1 v1 n) e2
11   -- Return the value of e2 and the combined trace
12   in (v2, t1 ++ t2)

```

Listing 7: Tracing let bindings

2.4 Array Tracing

Tracing on data structures like arrays provides us with a new problem that revolves around whether or not we wish to trace arrays away or not. We can see arrays as either structures that contain the data we are really after, which would require us to trace them away, or as data in their own right which we wish to keep in the trace. Both

scenarios provide us with interesting challenges.

Let us first talk about tracing arrays away. When we simply view arrays as another computational structure, they are not too complicated to trace away. When initializing an array, we just initialize all the individual values in the array. And when performing operations on items in the array, we instead perform those operations on the individual items again. That is, we do the operation like normal, but denote them as operations on separate items in the trace.

In Listing 8 we first add arrays and array operations. While we said earlier that constructors in the `TValue` ADT only needed strings for names if they are traced, we need to make an exception for arrays. This is because when working with arrays our expression will never refer to individual values in arrays, only to the array itself (and using its individual values from there). This means that to consistently refer to values that were in arrays in the original expression, we need to give a little more structure to the naming scheme. We do this by taking the name of the array, and adding the index of the item to create a name that is unique yet identifiable. Furthermore, we add in array operations: `iota` (or range operation) (`Iota`), indexing (`Idx`), sum (`Sum`), and map (`Map`). It should be noted that `iota` and indexing take an integer argument as part of their operator. For `iota` this is the size of the array to create, and for indexing this is naturally the index to get. While we could allow our language with integers (or by casting floats) to allow using in-language numbers, this really is not all that interesting. If those arguments were part of our trace, it would just mean tracing them like any other number.

```

1 data Value = ... | VArray [Float]
2
3 data TValue = ... | TArray String [Float]
4
5 data Expression
6   = ...
7   | EOp0 Op0
8   | EOp1 Op1 Expression
9
10 data Op0 = Iota Int
11
12 data Op1 = Idx Int | Sum
13
14 data Op2 = ... | Map

```

Listing 8: Adding arrays

Now with arrays added to our language, we can actually trace them. This is done in Listing 9, where we again extend the trace function, leaving out any patterns that remain unchanged.

First off, when encountering a literal array, or creating one with the `iota` operator, we need to initialize every individual value. This is fairly simple, it just requires us to walk through the array and initialize every value like when we were initializing literal real values.

Indexing, in this mode, is equal to variable reference due to our naming scheme. This means then that we do not need to trace anything here.

```

1 trace :: TEnvironment -> Expression -> (TValue, Trace)
2 trace n (ELift v) =
3     case v of
4         -- Tracing for reals and Booleans remain unchanged
5         (VReal v) -> ...
6         (VBool v) -> ...
7         (VArray v) -> let s = getName
8             -- For traceArrayLift see Listing 10
9             in (TArray s v, traceArrayLift s v 0)
10
11 -- With only iota, we could write this a little more curtly, but for clarity we leave it like this
12 trace n (EOp0 op) =
13     case op of
14         (Iota r) ->
15             let s = getName
16                 -- Define an array of size r, then lift is using traceArrayLift again
17                 v = [0.0 .. (r - 1)]
18                 -- For traceArrayLift see Listing 10
19             in (TArray s v, traceArrayLift s v 0)
20
21 trace n (EOp1 op e1) =
22     -- Again we first trace e1, and we get a name ready as well
23     let (v1, t1) = trace n e1
24         s = getName
25     in case (op, v1) of
26         -- Indexing is like variable reference, we do not need to add to the trace,
27         -- but we need to create the name to be consistent
28         (Idx i, TArray s1 v) =
29             -- Get the actual item using indexing (!!)
30             let x = v !! i
31                 s' = s1 ++ '!' : show i
32             in (TReal s' x, t1)
33         -- For traceArraySum see Listing 11
34         (Sum, TArray s1 v) =
35             let (vs, ts) = traceArraySum s1 v 0
36                 -- We must not forget to add the trace of e1 to our trace here
37             in (vs, ts ++ t1)
38
39 trace n (EOp2 op e1 e2) =
40     -- Again we first trace e1 and e2, and we get a name ready as well
41     let (v1, t1) = trace n e1
42         (v2, t2) = trace n e2
43         s = getName
44     in case (op, v1, v2) of
45         ...
46         (Map, TFunc f, TArray sa va) ->
47             -- For traceArrayMap see Listing 10
48             let (vm, tm) = traceArrayMap f sa va 0
49             -- Combine the traces
50             in (vm, tm ++ t1 ++ t2)

```

Listing 9: Tracing away arrays

```

1  -- traceArrayLift takes the name of the array, the contents, and the current index
2  traceArrayLift :: String -> [Float] -> Int -> Trace
3  -- Empty lists get no trace
4  traceArrayLift _ [] _ = []
5  traceArrayLift s (x:xs) i =
6      -- Create the name for this item from the array's name and the current index
7      let s' = s ++ '!' : show i
8      -- Trace x as a single real number
9      tx = TLift (TReal s' x)
10     -- Trace the rest of the array
11     txs = traceArrayLift s xs (i + 1)
12     -- Return the combined trace
13     in tx : txs
14
15  -- traceArrayMap takes the function to map, the name of the old array, the name of the new array,
16  -- the contents of the old array, and the current index
17  traceArrayMap :: (TValue -> (TValue, Trace)) -> String -> String -> [Float]
18  --> Int -> (TValue, Trace)
19  traceArrayMap _ _ sn [] _ = (TArray sn [], [])
20
21  traceArrayMap f so sn (x:xs) i =
22      -- Get the current value from the array with the right name
23      let current = TReal (so ++ '!' : show i) x
24      -- Get the result from the function application
25      (fv, ft) = f current
26      -- Get the results from the rest of the array
27      (xsv, xst) = traceArraymap f so sn xs (i + 1)
28      -- To add to the TArray and to rename fv we use this case-of statement
29      in case (fv, xsv) of
30          (TReal s' v, TArray _ xsv') ->
31              -- Add this item to the new array
32              let vn = TArray sn (v : xsv')
33              -- Rename fv in the function trace to the correct name
34              ft' = rename s' (sn ++ '!' : show i) ft
35              -- Finally return the new array and the combined trace
36              in (vn, ft' ++ xst)
37
38  rename :: String -> String -> Trace -> Trace

```

Listing 10: Tracing array instantiation and array mapping

The sum operator is a little more in-depth, as shown in Listing 11. However, this is a lot of code for a very simple principle, and a couple edge cases. The principle is, add the first two values in the array together, and then every following item to that result and so on. And we have edge cases for singleton and empty arrays. It is worth explicitly stating that every addition done whilst summing the array gets its own unique name and step in the trace. This means that an operation that is single step the original program, explodes to a bunch of steps (the length of the array minus one) in the trace. This is because we decided to trace away arrays, and we will see later on how we save ourselves from this by not tracing away arrays.

The map operator is funky in a way similar to sum. In essence, we take each item in the array and apply it to the function as expected. However, we run into a little problem with our naming scheme. For map, the items, once mapped on, are placed back into a new array. This means that, according to the scheme we laid out, the items in this array should be named in reference to the new array, however this is not something the call to trace in the function body considers. Luckily we can resolve this by renaming the returned value from that function, and changing the name in the trace. The signature of a function that does this is also included at the end of Listing 10, but its exact implementation is not of importance here.

While the concepts behind tracing away arrays are hopefully not too difficult to understand, it should be obvious from Listings 10 and 11 that the implementation becomes more complex. Now while that is not really a problem, we should really note that the trace becomes messier as well. This is especially problematic if we actually want to read the trace to see what is going on: not impossible, but also not pleasant, especially with large arrays. So perhaps we are tempted to keep arrays in the trace instead, or perhaps we are interested in the trace of arrays specifically.

Luckily for us, in large parts tracing while keeping arrays is fairly easy. This is because we can treat most operations like how we treated operations for real numbers. This has been done in Listing 12, except for map, where replace the tracing patterns from Listing 9.

In our current language, the main point of difficulty and interest is the map operation. Map takes in a function, which is not a type we wish to keep in our trace, however it produces an array which we wish to keep in our trace. While we might be tempted to just discard the function component, we cannot do that because it provides the trace from the original array to the new array. Without that information our trace is no longer a functional (straight-line) program.

The basic way to solve this, the naïve method, would be to attach an array of traces to the map operator, so they can be followed to derive the correct results. To easily do this we extend our **Traced** ADT with a special map constructor (**TMap**). We show this in Listing 13. The traces in the **TMap** constructor correspond with the application of the function to be mapped to the individual item, for each item. The string references the array the map is performed on.

Now, while the naïve way is fine in functionality, it does again create some overhead (a trace for each item in the array) by splitting the trace into multiple smaller traces. And if the function is the same for every item in the array, we may find ourselves saving a lot of redundant data. Now, this may be necessary: at the time we map a function over an array, we do not know if it will act the same for every input. Perhaps there is some control flow in the function body that checks if a number is even, or a factor


```

1  -- traceArraySum starts the trace, and traceArraySum' completes it
2  -- This is necessary because we do not know the number of items in the array
3  -- traceArraySum takes only the array to sum
4  traceArraySum (TArray _ []) =
5      let s = getName
6          v = TReal s 0
7          -- The sum of an empty array means just lifting the value 0
8          in (v, [(s, TLift v)])
9
10 traceArraySum (TArray _ [x]) =
11     let s = getName
12         v = TReal s x
13         -- The sum of a singleton array is just that one value
14         in (v, [(s, TLift v)])
15
16 traceArraySum (TArray sa (x:y:z)) =
17     -- When summing on a larger array, the first sum is of the first two items
18     let sx = sa ++ '!0'
19         sy = sa ++ '!1'
20         s = getName
21         v = TReal s (x + y)
22         -- Get the result, and the trace of the rest of the array with traceArraySum'
23         (rv, rt) = traceArraySum' (TArray sa z) 2 v
24         -- Return the final result, but do not forget the trace of the first sum
25         in (rv, (s, TOp2 Add sx sy) : rt)
26
27 -- traceArraySum' takes the array we sum over, the current index, and the last calculated value
28 traceArraySum' :: TValue -> Int -> TValue -> (TValue, Trace)
29 -- When we are done, return the value
30 traceArraySum' (TArray _ []) _ v = (v, [])
31
32 traceArraySum' (TArray sa (x:xs)) i (TReal sr r) =
33     -- Get the name for this item
34     let sx = sa ++ '!' : show i
35         -- Get the name for this addition step
36         s = getName
37         -- Get the result of the rest of the array
38         (v, t) = traceArraySum' (TArray sa xs) i (TReal s (x + r))
39         -- Return the final result, and add this steps addition to the trace
40         in (v, (s, TOp2 Add sx sr) : t)

```

Listing 11: Tracing the sum operator

```

1 trace :: TEnvironment -> Expression -> (TValue, Trace)
2 trace n (ELift v) =
3     case v of
4         (VReal v) -> ...
5         (VBool v) -> ...
6         (VArray v) ->
7             let s = getName
8                 -- Literal lifting of arrays becomes real simple
9             in (TArray s v, [(s, TLift (TArray s v))])
10
11 trace n (EOp0 op) =
12     case op of
13         (Iota r) ->
14             let s = getName
15                 v = [0.0 .. (r - 1)]
16                 -- Iota again becomes very similar to literal array lifting
17             in (TArray s v, [(s, TLift (TArray s v))])
18
19 trace n (EOp1 op e1) =
20     -- We trace e1 first, and create a name just in case
21     let (v1, t1) = trace n e1
22         s = getName
23     in case (op, v1) of
24         (Idx i, TArray s1 v) =
25             let x = v !! i
26                 s' = s1 ++ '!' : show i
27                 -- Now we trace arrays, indexing becomes more relevant to add to our trace,
28                 -- as the individual item has not been defined before
29             in (TReal s' x, (s', TOp1 op s1) : t1)
30     -- Sum becomes very simple, just apply it to the array
31     (Sum, TArray s1 v) = (TReal s (sum v), (s, TOp1 Sum s1) : t1)

```

Listing 12: Tracing whilst keeping arrays

```

1 data Traced
2   = ...
3   | TMap [Trace] String
4
5 trace :: TEnvironment -> Expression -> (TValue, Trace)
6 trace n (EOp2 op e1 e2) =
7   -- We first trace e1 and e2, and generate a name
8   let (v1, t1) = trace n e1
9       (v2, t2) = trace n e2
10      s = getName
11  in case (op, v1, v2) of
12    ...
13    (Map, TFunc f, TArray sa va) ->
14      let (vs, ts) = traceMapNaive f sa s va 0
15          -- The trace becomes TMap, the collection of traces ts, on the old array v2 (with
16          -- name sa)
17      in (vs, [(s, TMap ts sa)])
18
19 -- traceMapNaive takes in the function to be mapped, the name of the old array, the name of the
20 -- new array, the contents of the old array, and the current index
21 traceMapNaive :: (TValue -> (TValue, Trace)) -> String -> String -> [Float]
22               -> Int -> (TValue, [Trace])
23 -- A map over an empty array returns the empty array and no traces
24 traceMapNaive _ _ sn [] _ = (TArray sn [], [])
25
26 traceMapNaive f so sn (x:xs) i =
27   -- Create specific names for the old and new value
28   let old = so ++ '!' : show i
29       new = sn ++ '!' : show i
30       -- Apply the function, getting the value for x and its trace
31       (xv, xt) = f (TReal old x)
32       -- Apply the function for the rest of the map
33       (xsv, xst) = traceMapNaive f so sn xs (i + 1)
34       -- We use a case-of statement to append xv to xsv and to rename xv in xt
35   in case (xv, xsv) of
36     (TReal s' v, TArray _ vs) ->
37       let xt' = rename s' new xt
38       in (TArray sn (v : vs), xt' : xst)

```

Listing 13: Tracing map while keeping arrays, naïvely

of three, or something else entirely. In such a case, having a trace for each item may be strictly necessary. However, it also highlights for which functions it may not be: functions without control flow or branching. After all, these functions are little straight-line programs, and should act the same no matter on what input they are applied. Writing a function that checks if a the body of a lambda abstraction contains branching is very simple for this language: currently the only expression term that can introduce branching is the if-then-else statement. Unfortunately, we cannot check that at the moment the we trace a map operator. This is because any function here would already have been abstracted to a `TArray` value. So, we would need to check for branching when we are abstracting the function and we also need a way to convey if a specific instance of `TFunc` contains branching or not. We write branch-checking into functions in Listing 14. For most terms we can just commute the branch checking to the arguments of that term, but there are a couple exceptions. If-then-else statements are the definition of branching in our language, so they return ‘true’, and no branching can occur in literal instantiation (`ELift`) or nullary operators (`Op0`) (literal instantiation can also be rewritten as a nullary operator), so they always return ‘false’. Only for variable reference, which may return a value without actually providing a code to check, we need to see if the value is a function, and whether it has the branching flag set or not. This works because we set the branching flag when functions are defined using abstraction, and because functions may not be entered as literals.

```

1 data TValue
2   = ...
3   -- Add a branching flag to TFunc
4   | TFunc Bool (TValue -> (TValue, Trace))
5
6 branchCheck :: TEnvironment -> Expression -> Bool
7 -- Encountering an if-else-statement means a encountering a branch
8 branchCheck _ (EIf _ _ _) = True
9
10 branchCheck n (EApply e1 e2) = branchCheck n e1 || branchCheck n e2
11 branchCheck n (ELambda _ e1) = branchCheck n e1
12 branchCheck n (ELet _ e1 e2) = branchCheck n e1 || branchCheck n e2
13 -- ELift is always false, because lifting functions is not allowed
14 branchCheck _ (ELift) = False
15 branchCheck _ (EOp0 _) = False
16 branchCheck n (EOp1 _ e1) = branchCheck n e1
17 branchCheck n (EOp2 _ e1 e2) = branchCheck n e1 || branchCheck n e2
18
19 -- If our variable contains a function we need to check what it has the branching flag set to
20 branchCheck n (ERef s1) = case n ! s1 of
21   (TFunc b _) -> b
22   _           -> False

```

Listing 14: Checking for branches

With our branch checking defined we still need to talk about how we actually apply that and make a trace for map that requires less information. The basic idea here is that we can essentially perform vectorization of our function on the array in our trace: we rewrite the trace such that the function is “applied” to the whole array, rather than its individual items. Now without support for this in our language, this basically amounts to syntactic sugar in our trace, however it will provide us with a much clearer trace. This has been done in Listing 15, where we again add a map operator to our `Traced` ADT. This is because we may need to use the naïve method if a function contains branching

and we cannot vectorize it. In Listing 15 we still use `traceMapNaive` to actually map over our array. This is because we need to get the value of the array regardless, and our function value (`TFunc`) will return traces regardless if we need them or not. Then we can just take the first trace returned by the naïve map tracing, and rename all references to the first item of both the new and old arrays, to references of the whole old and new arrays respectively. For this end we define a function `deepRename` at the end of Listing 15. Like with the renaming function in Listing 10, the implementation of this function is not all that interesting: since all a `Trace` object is, is a list of tuples with a name that may need renaming and a `Traced` constructor referencing zero to two strings that may need renaming. All `deepRename` would do is go over these items and rename any occurrences it finds.

```

1 data Traced
2   = ...
3   -- We leave the naive TMap untouched
4   | TMap [Trace] String
5   -- And add a new one for vectorized traces
6   | TMapV Trace  String
7
8 trace :: TEnvironment -> Expression -> (TValue, Trace)
9 trace n (ELambda s1 e1) =
10   -- We add branch checking when we handle abstraction
11   let b = branchCheck e1
12       f = TFunc b (\x -> trace (insert s x) e1)
13   in (f, [])
14
15 trace n (EOp2 op e1 e2) =
16   let (v1, t1) = trace n e1
17       (v2, t2) = trace n e2
18       s = getName
19   in case (op, v1, v2) of
20     ...
21     (Map, TFunc b f, TArray sa va) ->
22       -- We first get the result array (and all the traces) using the naive method
23       let (vs, ts) = traceMapNaive f sa s va 0
24       in if b
25         -- If the function contains branching, use the naive method
26         then (vs, (s, TMap ts sa) : t1 ++ t2)
27         -- Otherwise use the new method
28         else let t = vectorizeTrace sa s (head ts)
29              in (vs, (s, TMapV t sa) : t1 ++ t2)
30
31 vectorizeTrace :: String -> String -> Trace -> Trace
32 -- Rename the references to individual items to the whole array
33 vectorizeTrace so sn t = deepRename iso so (deepRename isn sn t)
34 -- The names for the individual items in this trace
35 where iso = so ++ '!0'
36       isn = sn ++ '!0'
37
38 deepRename :: String -> String -> Trace -> Trace

```

Listing 15: Array mapping with trace vectorization

What is important to take away from the shenanigans with the map operators is that, whilst our definitions and correctness assertions from Sections 2 and 2.1 gave us some

guidance, there is ultimately no single way to trace everything. The most important factor here is to keep reminding ourselves of the information we wish to keep in the trace. Not only the value types, but we also need the information needed to actually run the trace as a program. Keeping this in mind, it becomes much more obvious how to trace the map operation.

3 Automatic Differentiation

Hello, world!

References

- [1] Charles C Margossian. A review of automatic differentiation and its efficient implementation. *Wiley interdisciplinary reviews: data mining and knowledge discovery*, 9(4):e1305, 2019.
- [2] Louis B Rall. The arithmetic of differentiation. *Mathematics Magazine*, 59(5):275–282, 1986.
- [3] Seppo Linnainmaa. Taylor expansion of the accumulated rounding error. *BIT Numerical Mathematics*, 16(2):146–160, 1976.
- [4] Andreas Griewank. Achieving logarithmic growth of temporal and spatial complexity in reverse automatic differentiation. *Optimization Methods and software*, 1(1):35–54, 1992.
- [5] Christian H Bischof and H Martin Bücker. Computing derivatives of computer programs. Technical report, Argonne National Lab., IL (US), 2000.
- [6] Claus Bendtsen and Ole Stauning. Fadbad, a flexible c++ package for automatic differentiation. Technical report, Technical Report IMM-REP-1996-17, Department of Mathematical Modelling . . . , 1996.
- [7] Andreas Griewank, David Juedes, and Jean Utke. Algorithm 755: Adol-c: A package for the automatic differentiation of algorithms written in c/c++. *ACM Transactions on Mathematical Software (TOMS)*, 22(2):131–167, 1996.
- [8] Pierre Aubert, Nicolas Di Césaré, and Olivier Pironneau. Automatic differentiation in c++ using expression templates and. application to a flow control problem. *Computing and Visualization in Science*, 3(4):197–208, 2001.
- [9] Robin J Hogan. Fast reverse-mode automatic differentiation using expression templates in c++. *ACM Transactions on Mathematical Software (TOMS)*, 40(4):1–16, 2014.
- [10] Laurent Hascoet and Valérie Pascual. The tapenade automatic differentiation tool: Principles, model, and specification. *ACM Transactions on Mathematical Software (TOMS)*, 39(3):1–43, 2013.
- [11] Andreas Griewank and Andrea Walther. *Evaluating derivatives: principles and techniques of algorithmic differentiation*. SIAM, 2008.
- [12] Fei Wang, Daniel Zheng, James Decker, Xilun Wu, Grégory M Essertel, and Tiark Rompf. Demystifying differentiable programming: Shift/reset the penultimate backpropagator. *Proceedings of the ACM on Programming Languages*, 3(ICFP):1–31, 2019.
- [13] Faustyna Krawiec, Simon Peyton Jones, Neel Krishnaswami, Tom Ellis, Richard A Eisenberg, and Andrew Fitzgibbon. Provably correct, asymptotically efficient, higher-order reverse-mode automatic differentiation. *Proceedings of the ACM on Programming Languages*, 6(POPL):1–30, 2022.
- [14] Matthijs Vákár and Tom Smeding. Chad: Combinatory homomorphic automatic differentiation. *ACM Transactions on Programming Languages and Systems (TOPLAS)*, 44(3):1–49, 2022.

- [15] Conal Elliott. The simple essence of automatic differentiation. *Proceedings of the ACM on Programming Languages*, 2(ICFP):1–29, 2018.
- [16] Robert Schenck, Ola Rønning, Troels Henriksen, and Cosmin E Oancea. Ad for an array language with nested parallelism. In *SC22: International Conference for High Performance Computing, Networking, Storage and Analysis*, pages 1–15. IEEE, 2022.
- [17] Tom J Smeding and Matthijs IL Vákár. Efficient dual-numbers reverse ad via well-known program transformations. *Proceedings of the ACM on Programming Languages*, 7(POPL):1573–1600, 2023.
- [18] Christian Bischof. *Issues in parallel automatic differentiation*. Argonne National Lab., 1991.
- [19] Christian Bischof, Niels Guertler, Andreas Kowarz, and Andrea Walther. Parallel reverse mode automatic differentiation for openmp programs with adol-c. In *Advances in automatic differentiation*, pages 163–173. Springer, 2008.
- [20] Dougal Maclaurin. *Modeling, inference and optimization with composable differentiable procedures*. PhD thesis, Harvard University, Graduate School of Arts & Sciences, 2016.
- [21] Martín Abadi, Ashish Agarwal, Paul Barham, Eugene Brevdo, Zhifeng Chen, Craig Citro, Greg S Corrado, Andy Davis, Jeffrey Dean, Matthieu Devin, et al. Tensorflow: Large-scale machine learning on heterogeneous distributed systems. *arXiv preprint arXiv:1603.04467*, 2016.
- [22] Roy Frostig, Matthew James Johnson, and Chris Leary. Compiling machine learning programs via high-level tracing. *Systems for Machine Learning*, 4(9), 2018.
- [23] Bo Joel Svensson and Josef Svenningsson. Defunctionalizing push arrays. In *Proceedings of the 3rd ACM SIGPLAN Workshop on Functional High-performance Computing*, pages 43–52, 2014.
- [24] Guy E Blelloch. *NESL: a nested data parallel language*. Carnegie Mellon Univ., 1992.
- [25] Shail Aditya, Arvind, Jan-Willem Maessen, Lennart Augustsson, and Rishiyur S Nikhil. Semantics of ph: A parallel dialect of haskell. In *In Proceedings from the Haskell Workshop (at FPCA 95)*, pages 35–49, 1995.
- [26] Jonathan Michael David Hill. *Data parallel lazy function programming*. PhD thesis, Queen Mary, University of London, 1995.
- [27] Christoph A Herrmann and Christian Lengauer. Parallelization of divide-and-conquer by translation to nested loops. *Journal of functional programming*, 9(3):279–310, 1999.
- [28] Nils Ellmenreich, Christian Lengauer, and Martin Griebl. Application of the polytope model to functional programs. In *Languages and Compilers for Parallel Computing: 12th International Workshop, LCPC’99 La Jolla, CA, USA, August 4–6, 1999 Proceedings 12*. Springer, 2000.
- [29] Manuel MT Chakravarty, Gabriele Keller, Roman Lechtchinsky, and Wolf Pfannenstiel. Nepal—nested data parallelism in haskell. In *Euro-Par 2001 Parallel Processing: 7th International Euro-Par Conference Manchester, UK, August 28–31, 2001 Proceedings 7*, pages 524–534. Springer, 2001.

- [30] Gabriele Cornelia Keller. *Transformation-based implementation of nested data parallelism for distributed memory machines*. PhD thesis, Technical University of Berlin, Germany, 1999.
- [31] Gabriele Keller and Manuel MT Chakravarty. Flattening trees. In *Euro-Par'98 Parallel Processing: 4th International Euro-Par Conference Southampton, UK, September 1–4, 1998 Proceedings 4*, pages 709–719. Springer, 1998.
- [32] Manuel MT Chakravarty and Gabriele Keller. More types for nested data parallel programming. *ACM SIGPLAN Notices*, 35(9):94–105, 2000.
- [33] Manuel MT Chakravarty and Gabriele Keller. Functional array fusion. In *Proceedings of the Sixth ACM SIGPLAN International Conference on Functional Programming*, pages 205–216, 2001.
- [34] Dror G Feitelson. Packing schemes for gang scheduling. In *Job Scheduling Strategies for Parallel Processing: IPPS'96 Workshop Honolulu, Hawaii, April 16, 1996 Proceedings 2*, pages 89–110. Springer, 1996.
- [35] Gabriele Keller and Manuel MT Chakravarty. On the distributed implementation of aggregate data structures by program transformation. In *Parallel and Distributed Processing: 11th IPPS/SPDP'99 Workshops Held in Conjunction with the 13th International Parallel Processing Symposium and 10th Symposium on Parallel and Distributed Processing San Juan, Puerto Rico, USA, April 12–16, 1999 Proceedings 13*, pages 108–122. Springer, 1999.
- [36] Manuel MT Chakravarty, Roman Leshchinskiy, Simon Peyton Jones, Gabriele Keller, and Simon Marlow. Data parallel haskell: a status report. In *Proceedings of the 2007 workshop on Declarative aspects of multicore programming*, pages 10–18, 2007.
- [37] Simon Peyton Jones, Roman Leshchinskiy, Gabriele Keller, and Manuel MT Chakravarty. Harnessing the multicores: Nested data parallelism in haskell. In *IARCS Annual Conference on Foundations of Software Technology and Theoretical Computer Science*. Schloss Dagstuhl-Leibniz-Zentrum für Informatik, 2008.
- [38] Sven-Bodo Scholz. Single assignment c-functional programming using imperative style. In *Proceedings of IFL*, volume 94, 1994.
- [39] Sven-Bodo Scholz. Single assignment c: efficient support for high-level array operations in a functional setting. *Journal of functional programming*, 13(6):1005–1059, 2003.
- [40] Clemens Grelck and Sven-Bodo Scholz. Generic parallel array programming in sac. In *21. Workshop der GI-Fachgruppe 2.1.4 Programmiersprachen und Rechenkonzepte, Bad Honnef, Germany*, pages 43–53, 2005.
- [41] Gabriele Keller, Manuel MT Chakravarty, Roman Leshchinskiy, Simon Peyton Jones, and Ben Lippmeier. Regular, shape-polymorphic, parallel arrays in haskell. *ACM Sigplan Notices*, 45(9):261–272, 2010.
- [42] C Barry Jay and J Robin B Cockett. Shapely types and shape polymorphism. In *Programming Languages and Systems—ESOP'94: 5th European Symposium on Programming Edinburg, UK, April 11–13, 1994 Proceedings 5*, pages 302–316. Springer, 1994.

- [43] Manuel MT Chakravarty, Gabriele Keller, Sean Lee, Trevor L McDonell, and Vinod Grover. Accelerating haskell array codes with multicore gpus. In *Proceedings of the sixth workshop on Declarative aspects of multicore programming*, pages 3–14, 2011.
- [44] Justin Slepak, Olin Shivers, and Panagiotis Manolios. An array-oriented language with static rank polymorphism. In *ESOP*, volume 8410, pages 27–46. Springer, 2014.
- [45] Kenneth E Iverson. A programming language. In *Proceedings of the May 1-3, 1962, spring joint computer conference*, pages 345–351, 1962.
- [46] Troels Henriksen, Niels GW Serup, Martin Elsmann, Fritz Henglein, and Cosmin E Oancea. Futhark: purely functional gpu-programming with nested parallelism and in-place array updates. In *Proceedings of the 38th ACM SIGPLAN Conference on Programming Language Design and Implementation*, pages 556–571, 2017.
- [47] Dougal Maclaurin, Alexey Radul, Matthew J Johnson, and Dimitrios Vytiniotis. Dex: array programming with typed indices. In *Program Transformations for ML Workshop at NeurIPS 2019*, 2019.
- [48] Adam Paszke, Daniel Johnson, David Duvenaud, Dimitrios Vytiniotis, Alexey Radul, Matthew Johnson, Jonathan Ragan-Kelley, and Dougal Maclaurin. Getting to the point. index sets and parallelism-preserving autodiff for pointful array programming. *arXiv preprint arXiv:2104.05372*, 2021.

A ADT Evaluation

In Section 2.2, we introduced an extended lambda calculus. In this section we will quickly go over how an evaluator function for this ADT would look like in Haskell. We define our evaluator function in Listing 16, using the definitions of `Expression`, `Value`, and `Environment` from Listing 3.

```

1 eval :: Environment -> Expression -> Value
2
3 eval n (EApply e1 e2) =
4     -- Evaluate e1 and e2 first
5     let v1 = eval n e1
6         v2 = eval n e2
7     in case v1 of
8         -- Only apply v1 to v2 if v1 is a function as expected
9         VFunc f -> f v2
10        _       -> error "Type mismatch in eval/EApply"
11
12 eval n (EIf e1 e2 e3) =
13     -- Evaluate e1 as the condition of the if-then-else statement
14     case eval n e1 of
15         -- If e1 evaluates to true, evaluate e2
16         VBool True -> eval n e2
17         -- Otherwise, evaluate e3
18         VBool False -> eval n e3
19         _           -> error "Type mismatch in eval/EIf"
20
21 -- For abstractions, we return the function by moving the evaluation into the body.
22 -- Where we insert the anonymous value x into the environment as it was when the
23 -- function was defined.
24 eval n (ELambda s1 e1) = VFunc $ \x -> eval (insert s1 x n) e1
25
26 eval n (ELift v1) = v1
27
28 eval n (EOp2 op e1 e2) =
29     -- Evaluate e1 and e2 first
30     let v1 = eval n e1
31         v2 = eval n e2
32     in case (op, v1, v2) of
33         -- This case syntax allows us to select for the right op with the right
34         -- value types at the same time.
35         (Add, VFloat a, VFloat b) -> VFloat $ a + b
36         (Equ, VBool a, VBool b) -> VBool $ a == b
37         (Equ, VFloat a, VFloat b) -> VBool $ a == b
38         (Mul, VFloat a, VFloat b) -> VFloat $ a * b
39         (Neq, VBool a, VBool b) -> VBool $ a /= b
40         (Neq, VFloat a, VFloat b) -> VBool $ a /= b
41         _ -> error "Type mismatch in eval/EOp2"
42
43 -- Resolving references means getting the value from the environment by name.
44 eval n (ERef s1) = n ! s1

```

Listing 16: ADT Evaluator